


Research Portal — Caseworker Guide


How to use the AI Support Agent admin portal as a **caseworker**: the day-to-day support role. You work from the **summary data tier** — AI session summaries, risk and crisis signals, screeners, moods, check-ins, and safety plans — but never verbatim transcripts; those require the therapist role. Screenshots come from a practice sandbox with synthetic clients; your account starts in one (the amber banner) so you can learn every workflow safely.

1. Triage — your landing page




Your day starts here: caseload size, who needs attention, open crisis flags, and open escalations, with the roster ranked by urgency. Chips tell you why someone surfaced — open crisis flag, low check-in mood, overdue practice, safety plan on file. Click a client to open their profile.

2. The client profile — your data tier in action



You see the session timeline (titles, durations, crisis tags), check-ins, safety plans, and the care-notes panel. The amber notice explains what you *don't* see: the AI clinical brief and unredacted clinical timeline entries are therapist-only. That's the design, not a bug — your view is built for support coordination, and anything needing clinical eyes gets escalated (section 4). Your case notes live here too: drafts autosave, and **signing a note makes it immutable**.


3. Work Queue — ack and resolve



Crisis flags, worsening screeners, unread-message alerts, and inactivity items for your clients land here. **Ack** claims an item so the team knows you're on it; **Resolve** closes it with a short note for the audit trail.

Set your notification preferences below the queue — keep "always email urgent safety items immediately" on.

4. Escalations — getting clinical help



When something needs clinical judgment — disclosed ideation, a worrying pattern, anything beyond support coordination — raise an escalation: pick the client, set urgency, and write what changed, what you observed, and what you need from the therapist. The therapist acknowledges and takes it from there; you can track status under *Raised by me*. **When in doubt, escalate.** Nobody has ever been faulted for asking for clinical eyes too early.

5. Also in your nav

- **Messages** — secure async messaging with your clients (start a thread from their profile). Replies are expected within 1-2 business days; every participant message is safety-scanned automatically, so the inbox is never the only safety net.
- **Crisis Management** — the crisis events and risk history for your caseload.
- **Adverse Events** — visibility into IRB adverse-event reports involving your clients.
- **Caseload** — your client list and one-time invite links for registering new clients.

6. What you can't do (and who can)

- Read transcripts, watch live sessions, or see analytics — therapist/researcher tier.
- Create accounts or change roles — researcher only.
- Change system configuration or study settings — researcher only.

Practice everything in your sandbox first — synthetic clients, nothing reaches real participants, pagers, or email. If you ever need support yourself: this is heavy work; the same 988 and CAPS resources on every screen are for staff too. Version 2026-09-04.