

Research Portal — Researcher Guide

How to use the AI Support Agent admin portal as a **researcher** — the study-operations role with the widest access. Screenshots come from the staging environment with synthetic data. Sign in at the study admin URL with your researcher account; enable MFA from the shield icon in the header on your first login.

1. Your home: Session History

AI Therapist Research & Therapist Audit Portal

Welcome, Tutorial_researcher (researcher)

OPERATIONS

Triage

Live Monitoring

Sessions

Analytics

SAFETY

Crisis Management

Adverse Events

PEOPLE

Users

Caseload

User Sessions

Rate Limits

RESEARCH

System Prompts

Knowledge Base

Evals

Redaction Review

Consent Versions

Session History

Search session ID, name, or username...

mm/dd/yyyy

mm/dd/yyyy

Min message

Max message


Advanced Filters

Clear All

User	Session Name	Session ID	Start Time	End Time	Ended By	Duration	Status	Messages	Actions
demo	"Managing Stress through Expression"	sess_EDv5qhttyu0mQobpUy679	Aug 17, 2026, 11:07 AM	Aug 17, 2026, 11:11 AM	User	4m 19s	ended	29 (70 / 11A)	View Edit Delete
demo	"Navigating grief and loss"	sess_EAJq1bxDHhbPKrrR0KQrF	Aug 7, 2026, 12:44 PM	Aug 7, 2026, 12:45 PM	User	1m 24s	ended	15 (40 / 6A)	View Edit Delete
test	"Mindful Check-In and Awareness"	sess_EB1ZZ1rRHQB2AHDHrt3YL	Aug 1, 2026, 10:58 PM	Aug 1, 2026, 11:05 PM	demo	6m 44s	ended	23 (80 / 11A)	View Edit Delete
demo	"Urgent safety concerns"	sess_EB8HzHTYPQ7iIfce1Pa85	Aug 1, 2026, 10:21 PM	Aug 1, 2026, 10:24 PM	User	3m 28s	ended	25 (40 / 8A)	View Edit Delete
demo	Please provide the details of the therapy session for me to summarize.	sess_E7rpAfcfj104db0rJx5fQ	Jul 31, 2026, 6:15 PM	Jul 31, 2026, 6:18 PM	User	2m 43s	ended	4 (80 / 1A)	View Edit Delete
test	Please provide the details of the therapy session for me to generate a title.	sess_E7rxvPhZVivVLbwD7ia3Hz	Jul 31, 2026, 6:07 PM	Jul 31, 2026, 6:12 PM	User	5m 0s	ended	21 (90 / 9A)	View Edit Delete
demo	It seems there was an error and I didn't receive the details of the therapy session to summarize. Please	sess_E7ZnLys1LaHRKgzDBInhX	Jul 30, 2026, 10:43 PM	Jul 30, 2026, 10:51 PM	demo	8m 25s	ended	48 (120 / 28A)	View Edit Delete
demo	Please provide the details of the therapy session for me to summarize.	sess_E7Z7s0C4RNoHfx1RYts4v	Jul 30, 2026, 10:27 PM	Jul 30, 2026, 10:34 PM	User	7m 31s	ended	38 (110 / 14A)	View Edit Delete
	Please provide the details of the therapy session you'd like	sess_E7Z7s0C4RNoHfx1RYts4v	Jul 22, 2026,	Jul 22, 2026,				12 (20 /	View Edit Delete


Every session in the study, searchable by session ID, name, or username, with date and message-count filters. **View** opens the full transcript and insights; researchers hold the full-content data tier, so you see verbatim (redacted) conversation text. The demo/test rows you see here are excluded from all research exports automatically.

2. Live Monitoring




Real-time view of active sessions with counts for active users, messages, and crisis sessions. During study hours this is the "is anything happening right now" screen; the Crisis Only toggle filters to flagged sessions. From a live session you can observe and, when warranted, use the sideband controls to inject a message or end a session.

3. Analytics




Four tabs — Usage, Performance, Cost, Tools — over any date range. Session activity, voice-vs-text share, average duration, model spend, and which therapeutic tools the agent actually invokes. This is descriptive operations data; the research dataset proper comes from Export (section 9).

4. Safety: Crisis Management and Adverse Events



Crisis Management aggregates every crisis event, intervention, and risk assessment, with tabs for the full histories. Flags show what triggered them (auto vs manual) and the risk factors involved.



Adverse Events is your IRB duty station: high-severity crisis flags and distress survey responses auto-draft AE reports here with review deadlines (the nav badge turns red when something is overdue). Review each draft, complete it, and sign off — this queue is the source for IRB reporting.

5. Users

AI Therapist Research & Therapist Audit Portal

Welcome, Tutorial_researcher (researcher)

OPERATIONS

Triage

Live Monitoring

Sessions

Analytics

SAFETY

Crisis Management

Adverse Events

PEOPLE

Users

Caseload

User Sessions

Rate Limits

RESEARCH

System Prompts

Knowledge Base

Evals

Redaction Review

Consent Versions

User Management

13 of 13 users

+ Add User

Search by username or user ID...

Filters

USER ID	USERNAME	ROLE	PREFERRED VOICE	LANGUAGE	MFA STATUS	MEMORY	RISK CONTEXT	CREATED AT	ACTIONS
76	tutorial_researcher	researcher	Cedar	en	Disabled	—	—	9/4/2026	View profile Edit Delete
74	demo_ffd8d386	demo	Cedar	en	N/A	—	—	9/4/2026	View profile Edit Delete
73	demo_538ffb3ac	demo	Cedar	en	N/A	—	—	9/4/2026	View profile Edit Delete
72	demo_0e540b34	demo	Cedar	en	N/A	—	—	9/4/2026	View profile Edit Delete
71	demo_45f83c9e	demo	Cedar	en	N/A	—	—	9/4/2026	View profile Edit Delete
70	demo_b77eaf94	demo	Cedar	en	N/A	—	—	9/4/2026	View profile Edit Delete
67	qualtrics-e2e-test	participant	Cedar	en	N/A	Off	Off	9/3/2026	View profile Edit Delete

The account directory. Researchers are the only role that can create accounts and change roles — this is where you add research assistants, therapists, and caseworkers. MFA status, memory setting, and risk-context sharing are all visible per user.

6. Study Ops

AI Therapist Research & Therapist Audit Portal

Welcome, Tutorial_researcher (researcher)

OPERATIONS

Triage

Live Monitoring

Sessions

Analytics

SAFETY

Crisis Management

Adverse Events

PEOPLE

Users

Caseload

User Sessions

Rate Limits

RESEARCH

System Prompts

Knowledge Base

Evals

Redaction Review

Consent Versions

Study Ops

Enrollment, arm balance, and protocol deviations (demo data excluded).

Run scan

ENROLLMENT

3 / 40

+0 anonymous sessions

ARM IMBALANCE

0.250

threshold 0.15

OPEN DEVIATIONS

0


0 major

SESSIONS / PARTICIPANT


1 at/above expected

2 below

Weekly enrollment



proactive_offering arm balance over threshold



Sessions per participant (reference: expected 4)



The protocol's control panel: enrollment against target, condition-arm balance with an imbalance alarm, sessions-per-participant against the expected count, and protocol deviations. **Run scan** checks for new deviations on demand; the same scan runs on a schedule. Demo data is excluded.

7. Qualtrics: Sync and Survey Data

AI Therapist Research & Therapist Audit Portal

Welcome, Tutorial_researcher (researcher)

OPERATIONS

Triage

Live Monitoring

Sessions

Analytics

SAFETY

Crisis Management

Adverse Events

PEOPLE

Users

Caseload

User Sessions

Rate Limits

RESEARCH

System Prompts

Knowledge Base

Evals

Redaction Review

Consent Versions

Qualtrics Sync

Scheduler

On

every 60 min

Last run

9/4/2026, 9:35:29 PM

trigger: scheduled

Surveys configured

5

baseline, weekly, exit, week12, withdrawal

Unlinked (finished)

7

responses with no participant

Linkage health

Survey	Total	Finished	Linked	Unlinked finished	Latest response
baseline	6	6	1	5	9/4/2026, 3:43:16 PM
exit	1	1	0	1	9/2/2026, 11:23:00 AM
weekly	4	4	3	1	9/4/2026, 3:45:01 PM

Last run results

Survey	Fetched	Upserted	Linked	Error
baseline	6	6	1	
weekly	4	4	3	
exit	1	1	0	
week12	0	0	0	
withdrawal	0	0	0	

Sync now

All five study surveys (baseline, weekly, exit, week-12, withdrawal) sync on a schedule; **Sync now** forces a run. Watch *Linkage health*: "unlinked finished" responses are ones that could not be matched to a participant (usually a mistyped study ID) and need human attention. Survey Data (next nav item) shows the per-participant completion matrix and scored instruments.

8. Evals

AI Therapist Research & Therapist Audit Portal

Welcome, Tutorial_researcher (researcher)

OPERATIONS

Triage

Live Monitoring

Sessions

Analytics

SAFETY

Crisis Management

Adverse Events

PEOPLE

Users

Caseload

User Sessions

Rate Limits

RESEARCH

System Prompts

Knowledge Base

Evals

Redaction Review

Consent Versions

Evals

Simulation Runs

Run now

voice variations 1 Start run

Voice runs open real Realtime sessions (a few minutes, real API spend) and leave playable recordings. Replay re-drives recent redacted sessions through the current pipeline.

Nightly run

☒ enabled voice at 9:00 UTC variations 1 Save


Runs the voice suite nightly at 09:00 UTC (2 AM local).

RUN	STARTED	SUITE	TRIGGER	PASSED	COST	COMMIT
#24	9/4/2026, 3:03:53 AM	voice	nightly	2/2	\$0.0019	-
#23	9/3/2026, 3:02:12 AM	voice	nightly	1/2	\$0.0019	-
#22	9/2/2026, 3:03:12 AM	voice	nightly	2/2	\$0.0019	-
#21	9/1/2026, 3:03:10 AM	voice	nightly	1/2	\$0.0018	-
#20	8/31/2026, 3:03:09 AM	voice	nightly	2/2	\$0.0018	-
#19	8/30/2026, 3:03:08 AM	voice	nightly	1/2	\$0.0018	-
#18	8/29/2026, 3:03:06 AM	voice	nightly	2/2	\$0.0019	-
#17	8/28/2026, 3:03:05 AM	voice	nightly	2/2	\$0.0018	-

The safety-evaluation harness: nightly simulated voice sessions run through the real pipeline, judged against gating assertions, with pass-rate trends and per-run cost. Also home to eval calibration (LLM

judge vs human ratings) and drift alerts. If a nightly run fails, the drill-down shows exactly which assertion broke.

9. Export



The **de-identified research dataset**: pseudonymized CSVs plus a generated codebook, deterministic for a given as-of date, demo data excluded, no message content or pseudonym mapping ever included. Redacted transcripts are a separate opt-in artifact handled as identifiable data. The lower form does operational JSON/CSV exports.

10. Onboarding staff: Sandbox Invites

AI Therapist Research & Therapist Audit Portal

Welcome, Tutorial_researcher (researcher)

OPERATIONS

Triage

Live Monitoring

Sessions

Analytics

SAFETY

Crisis Management

Adverse Events

PEOPLE

Users

Caseload

User Sessions

Rate Limits

RESEARCH

System Prompts

Knowledge Base

Evals

Redaction Review

Consent Versions

Sandbox Invites

One-time links that create an isolated, pre-seeded demo environment (a fresh sandbox organization with 6-9 synthetic clients) for a therapist or caseworker. Sandbox data never reaches research exports, crisis paging, or email.

Mint a batch

Count

Dashboard role

Batch label (optional)

Mint links

Links expire after 30 days if unused. Each link can be used exactly once.

Batch history

Refresh

Label	Role	Created	Expires	Used
tutorial-capture	Caseworker	9/4/2026	10/4/2026	1 / 1
tutorial-capture	Therapist	9/4/2026	10/4/2026	1 / 1
test	Caseworker	8/27/2026	9/26/2026	1 / 1
test	Caseworker	8/27/2026	9/26/2026	0 / 10

Mint one-time links that create an isolated, pre-seeded practice environment (6-9 synthetic clients) for a new therapist or caseworker. Sandbox data never reaches research exports, crisis paging, or email — it is the safe way for new staff to learn the portal before touching real participants. Their tutorial (see the therapist and caseworker guides) assumes they start here.

11. System Prompts

AI Therapist Research & Therapist Audit Portal

Welcome, Tutorial_researcher (researcher)

OPERATIONS

Triage

Live Monitoring

Sessions

Analytics

SAFETY

Crisis Management

Adverse Events

PEOPLE

Users

Caseload

User Sessions

Rate Limits

RESEARCH

System Prompts

Knowledge Base

Evals

Redaction Review

Consent Versions

System Prompts

Configure the AI system prompts for different session types

Show Preview

Reset

Save Changes

<> Available Variables

Use `{{crisis_text}}` to insert the configured crisis contact information dynamically. This will be replaced with the current crisis hotline, phone, and text line when the prompt is used.

Realtime (Voice) Sessions

Chat-Only Sessions

Therapeutic Modality

CBT-Informed

Description

System prompt for realtime voice therapy sessions

System Prompt

5243 characters

Purpose & Scope

You are an AI **therapeutic assistant** for adults, providing **general emotional support and therapeutic conversation** only. Use empathy and evidence-based self-help (e.g., **CBT, DBT, mindfulness, journaling**) to help users cope with stress, anxiety, and common emotions. Make it clear: you **support and guide**, not **replace a human therapist**. **Once**, at the start of the session, make clear you are **not licensed** and that your help is **not a substitute for professional therapy/medical care** – then do NOT repeat this disclaimer in later replies unless the participant asks about your limits or requests diagnosis/medical advice. Encourage seeking a **licensed therapist for serious issues**. Stay within **support, coping, active listening, and psycho-education**—no clinical claims.

Boundaries & Limitations


Never diagnose, give medication, or legal advice. **Avoid** medical or legal topics; instead, offer **non-medication coping, self-care, lifestyle tips, relaxation, and gentle suggestions**. Do not suggest specific drugs/supplements or treatment plans. If asked for diagnosis or medical/legal advice, **politely decline** and clarify your non-professional status. Never misrepresent your credentials. Do not set up treatment plans or contracts or act as a human/professional; **focus on user's goals and autonomy**, using open-ended questions and suggestions.

Crisis Protocol

The instructions the AI runs on, in three tabs: **Realtime (Voice) Sessions**, **Chat-Only Sessions**, and **Therapeutic Modality** (the CBT-informed layer that condition arms toggle). The `{{crisis_text}}` variable inserts the currently configured crisis contact line wherever it appears, so crisis numbers are edited once in System Config, never hardcoded here. **Show Preview** renders the final assembled prompt; **Save Changes** applies to *new* sessions only.

Study rule: during Phase 2 the prompt is a frozen instrument. The verbatim text is on file with the IRB, so a mid-study edit is a protocol event — it changes the intervention under test. Treat this page as read-and-verify: check the prompt matches the IRB-filed version, and coordinate with the PI (and an IRB amendment if substantive) before any edit. Every save is versioned and attributable.

12. Knowledge Base



The corpus the AI can retrieve mid-conversation: **psychoeducation** articles, **techniques**, and **worksheets**, each item tagged by topic and source. Only *approved* content is ever used in live sessions — new items sit in *Pending approval* until a researcher reviews them, so the approval gate is your quality control over what the AI can cite. Each card shows its source and license, when it was last retrieved, and Edit/Unapprove/Delete actions.

Two support tabs: **Test retrieval** lets you type a probe query ("can't sleep, mind racing") and see exactly which items the AI would pull, ranked — use it after adding content to confirm it surfaces. **Retrieval log** shows what was actually retrieved in real sessions, which tells you what participants are asking about and which items pull dead weight. Like the prompt, the corpus is part of the intervention: keep it stable during the study, and log any additions.

13. Redaction Review

AI Therapist Research & Therapist Audit Portal

Welcome, Tutorial_researcher (researcher)

OPERATIONS

Triage

Live Monitoring

Sessions

Analytics

SAFETY

Crisis Management

Adverse Events

PEOPLE

Users

Caseload

User Sessions

Rate Limits

RESEARCH

System Prompts

Knowledge Base

Evals

Redaction Review

Consent Versions

Redaction Verification

Review and edit redacted message content

Load More

assistant

text

Aug 28, 2026, 3:03 AM

ID: 991

Hi, I'm an AI mental health support assistant. I'm not a licensed therapist or doctor, and this isn't a substitute for professional care, but I can listen and offer coping ideas. If you're in crisis or feel unsafe, please contact 911 or your local emergency services.

You don't have to know what to say. We can start very simple: what's been on your mind lately, or what feeling shows up most—stress, worry, sadness, or just feeling stuck? If you'd like, you can also tell me what you hope this chat might help with, even if it's just "I want to feel a little less overwhelmed."

assistant

text

Aug 16, 2026, 9:31 AM

ID: 1527

That is a real win, silly-feeling or not. What did you notice when you contained the worrying to that window?

assistant

text

Aug 18, 2026, 7:44 AM

ID: 1614

So the catastrophe chain broke at step one. That gap between the prediction and the outcome is worth writing down – it is evidence your anxiety exaggerates.

assistant

text

Sep 3, 2026, 3:02 AM

ID: 1318

That "weight" feeling is really relatable, and it makes sense it would wear you down when you can't quite name it. You don't have to have a clear label for it to matter.

Every message is automatically de-identified (names, places, identifying details) and the original text is deleted after the retention window; the redacted version is what persists and what research exports may include. This page is the human check on that pipeline: it lists redacted messages, oldest-verified-last, and you spot-check that (a) no identifying detail slipped through and (b) the redaction didn't gut the clinical meaning. If something slipped through, edit the text right here — your correction replaces the stored redacted content. A periodic pass (for example weekly, and before any transcript export) is a protocol expectation, not busywork: this is the layer that makes the rest of the data pipeline safe.

14. Consent Versions

AI Therapist Research & Therapist Audit Portal

Welcome, Tutorial_researcher (researcher)

OPERATIONS

Triage

Live Monitoring

Sessions

Analytics

SAFETY

Crisis Management

Adverse Events

PEOPLE

Users

Caseload

User Sessions

Rate Limits

RESEARCH

System Prompts

Knowledge Base

Evals

Redaction Review

Consent Versions

Consent Versions

The versioned consent copy participants accept, per audience. The active version is the newest one whose effective date has passed. Each acceptance is stored with a hash of the exact text.

Publish new version

Research (IRB study)

IN EFFECT

Clinical (care team)

Shown to participants while deployment_mode is research.

2026-07-30.1

Active

Effective 7/29/2026, 6:00:00 PM · published by system-backfill · 133 acceptances

Preview

The versioned consent copy participants accept in-app, in two audiences: **Research (IRB study)** — shown while the deployment runs in research mode — and **Clinical (care team)**. The active version is the newest one whose effective date has passed; each participant's acceptance is stored with a hash of the exact text they saw, so you can always prove which words a given participant agreed to. **Preview** expands the full text; the acceptance count shows uptake.

Publish new version is how consent copy changes ship (for example after an IRB-approved amendment): publish with an effective date, and participants see the new text from then on — old acceptances keep pointing at the text they actually accepted. Never edit consent copy anywhere else; this page is the audit trail the IRB will ask for. In-app copy must always match the IRB-approved consent document version.

15. System Config and Data Retention

AI Therapist Research & Therapist Audit Portal

Welcome, Tutorial_researcher (researcher)

HFA

Logout

OPERATIONS

Triage

Live Monitoring

Sessions

Analytics

SAFETY

Crisis Management

Adverse Events

PEOPLE

Users

Caseload

User Sessions

Rate Limits

RESEARCH

System Prompts

Knowledge Base

Evals

Redaction Review

Consent Versions

System Configuration

Reset

Save Changes

Crisis Contact Information

This information is displayed to users in the AI's initial message and when crisis keywords are detected.

☐ Display crisis information to users

Hotline Name

988 Suicide & Crisis Lifeline

Phone Number

988

Crisis Text Line (optional)

Text HOME to 741741

On-Call Crisis Paging

When a session is flagged HIGH severity, an SMS pages this phone via the iMessage relay – in addition to the admin dashboard toast, which only reaches someone with the dashboard open. Failure to send never blocks the session.

☒ Page the on-call phone on high-severity crisis flags

On-Call Phone Number

+15551234567 (blank = use CRISIS_ALERT_PHONE env var)

Leave blank to fall back to the server's CRISIS_ALERT_PHONE environment variable.

System Config holds the study-level operational settings. The two that matter most: **Crisis Contact Information** — the hotline name, number, and text line that `{{crisis_text}}` injects into prompts and that participants see in crisis moments — and **On-Call Crisis Paging**, which SMS-pages the on-call phone whenever a session is flagged high severity (the dashboard toast only reaches someone with the dashboard open; the pager is the net that works at 2 AM). Verify the on-call number at every handoff. Changes here are study-level decisions — coordinate before touching them.

AI Therapist Research & Therapist Audit Portal

Welcome, Tutorial_researcher (researcher)

OPERATIONS

Triage

Live Monitoring

Sessions

Analytics

SAFETY

Crisis Management

Adverse Events

PEOPLE

Users

Caseload

User Sessions

Rate Limits

RESEARCH

System Prompts

Knowledge Base

Evals

Redaction Review

Consent Versions

Data Retention & Content Wipe

Automated deletion of original message content for IRB compliance

SCHEDULER

Active

Next: Sep 5, 2026, 9:00 PM

READY TO WIPE

0

messages eligible

AWAITING REDACTION

227

in redaction queue

WITH ERRORS

15

need attention

IRB Compliance

This system automatically deletes original message content after the retention period, leaving only the redacted (anonymized) content. This ensures PII is not retained longer than necessary for clinical purposes.

Retention Settings

Enable Automated Content Wipe

When enabled, original message content is automatically deleted on schedule

Retention Period (Hours)

24

24h

48h

72h

1 week

Original content is deleted after this many hours (only if redaction is complete)

Daily Wipe Time

03:00 AM

Automated wipe runs daily at this time (server timezone)

Data Retention runs the IRB data-minimization commitment: original (pre-redaction) message text is automatically deleted after the retention period — but only once redaction of that message is complete, so nothing is lost before its redacted copy exists. The status cards are your morning glance: *Ready to wipe* and *Awaiting redaction* are normal flow; ***With errors* counts messages whose redaction failed and therefore cannot be wiped — a nonzero number that persists needs attention, because those originals are outliving the window.** The retention period and daily wipe time are IRB-filed parameters; do not change them without a protocol conversation. Below the settings, the deletion audit log records every wipe.

16. Triage and Caseload: the care-team side

AI Therapist Research & Therapist Audit Portal

Welcome, Tutorial_researcher (researcher)

OPERATIONS

Triage

Live Monitoring

Sessions

Analytics

SAFETY

Crisis Management

Adverse Events

PEOPLE

Users

Caseload

User Sessions

Rate Limits

RESEARCH

System Prompts

Knowledge Base

Evals

Redaction Review

Consent Versions

Triage

Clients

0

Needs attention

0

Open crisis flags


0

Open escalations

0

Refresh

Two nav items connect you to the care-team side. **Triage** is the same attention-ranked view therapists work from — clients, who needs attention, open crisis flags, open escalations. For you it is an oversight surface: see what the care team sees, audit that flags and escalations are being worked, and pull care context when reviewing an adverse event. The counts are scoped to care assignments, so outside a care deployment (as in this staging shot) it sits at zero.



Caseload is where the researcher's people-management duty meets the care side: it assigns participants to therapists, and therapists only ever see their assigned clients. Create the therapist account in Users first, then make the assignment here. The care workflows themselves (working the queue, escalations, progress notes, participant messaging) belong to the therapist and caseworker roles — the companion guides cover them.

Questions or something looks wrong? The Report-a-problem button files directly to the engineering queue. Version 2026-09-04.